

Data Preparation

*Loading traning dataset

In [1]:

```
import pandas as pd
import numpy as np

data = pd.read_csv('LoansTrainingSet.csv')
data.head()
```

//anaconda/envs/python2/lib/python2.7/site-packages/IPython/core/interactiveshell.py:2717: DtypeWarning: Columns (16) have mixed types. Specify dtype option on import or set low_memory=False.

interactivity=interactivity, compiler=compiler, result=result)

Out[1]:

	Loan ID	Customer ID	Loan Status	Current Loan Amount	Term	Credit Score	Years in current job	Home Ownership	Annu Inco
0	000025bb-5694-4cff-b17d-192b1a98ba44	5ebc8bb1-5eb9-4404-b11b-a6eebc401a19	Fully Paid	11520	Short Term	741.0	10+ years	Home Mortgage	3369
1	00002c49-3a29-4bd4-8f67-c8f8fbc1048c	927b388d-2e01-423f-a8dc-f7e42d668f46	Fully Paid	3441	Short Term	734.0	4 years	Home Mortgage	4226
2	00002d89-27f3-409b-aa76-90834f359a65	defce609-c631-447d-aad6-1270615e89c4	Fully Paid	21029	Short Term	747.0	10+ years	Home Mortgage	9012
3	00005222-b4d8-45a4-ad8c-186057e24233	070bcecb-aae7-4485-a26a-e0403e7bb6c5	Fully Paid	18743	Short Term	747.0	10+ years	Own Home	3807
4	0000757f-a121-41ed-b17b-162e76647c1f	dde79588-12f0-4811-bab0-e2b07f633fcd	Fully Paid	11731	Short Term	746.0	4 years	Rent	5002



In [2]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 256984 entries, 0 to 256983
Data columns (total 19 columns):
Loan ID                256984 non-null object
Customer ID            256984 non-null object
Loan Status            256984 non-null object
Current Loan Amount    256984 non-null int64
Term                   256984 non-null object
Credit Score           195308 non-null float64
Years in current job    256984 non-null object
Home Ownership          256984 non-null object
Annual Income           195308 non-null float64
Purpose                 256984 non-null object
Monthly Debt            256984 non-null object
Years of Credit History 256984 non-null float64
Months since last delinquent 116601 non-null float64
Number of Open Accounts 256984 non-null int64
Number of Credit Problems 256984 non-null int64
Current Credit Balance  256984 non-null int64
Maximum Open Credit     256984 non-null object
Bankruptcies            256455 non-null float64
Tax Liens                256961 non-null float64
dtypes: float64(6), int64(4), object(9)
memory usage: 37.3+ MB
```

*As we can see from above, we have problems with some missing values & data types

*Let's take a look at "Current Loan Amount" Column

In [3]:

```
data['Current Loan Amount'].max(), data['Current Loan Amount'].min(),data['Current I
```

Out[3]:

```
(999999999, 505, 13713306.260237992)
```

*Selecting the rows of "Current Loan Amount" is lesser than 999999999

In [4]:

```
data = data[(data['Current Loan Amount'] < 999999999)]
```

In [5]:

```
data['Current Loan Amount'].describe()
```

Out[5]:

```
count      221774.000000
mean       13979.687389
std        8260.519207
min         505.000000
25%        7819.000000
50%       12058.000000
75%       19438.750000
max       41000.000000
Name: Current Loan Amount, dtype: float64
```

*Now, in "Current Loan Amount" column, the mean value is 13979.69 and median value is 12058.00

*Looking at unique values of "Years in current job", "Home Ownership", and "Purpose" columns

In [6]:

```
data['Years in current job'].unique()
```

Out[6]:

```
array(['10+ years', '4 years', '6 years', '5 years', 'n/a', '2 years',
      '< 1 year', '3 years', '1 year', '7 years', '9 years', '8 years'
      ], dtype=object)
```

In [7]:

```
data['Home Ownership'].unique()
```

Out[7]:

```
array(['Home Mortgage', 'Own Home', 'Rent', 'HaveMortgage'], dtype=object)
```

In [8]:

```
data['Purpose'].unique()
```

Out[8]:

```
array(['Debt Consolidation', 'other', 'Business Loan', 'Home Improvements',
      'Buy House', 'Other', 'Buy a Car', 'Medical Bills', 'Take a Trip',
      'Educational Expenses'], dtype=object)
```

Data Problems

Data Problems

1. Missing values: credit score (has nan), Annual Income, Months since last delinquent (has nan), Bankruptcies (has nan), Tax Liens (has nan)
2. Spelling differences & punctuation format Years in current job: delete years, 10+, <1, n/a;
Home Ownership: Home Mortgage == HaveMortgage; Purpose: Other == other; Monthly Debt: delete \$;
Maximum Open Credit: why is it Object? ; Years of Credit History: decimal 1 place
3. Duplicates rows
4. Some of Credit scores are too high, have to divide by 10
5. Current Loan Amount: too high 9999999 compared to their annual income
6. Data types: convert Credit Score - float64 to int64; convert Monthly Debt, Maximum Open Credit - object to float64


1. missing values: credit score (has nan)

In [9]:

```
credit_score_nan = data['Credit Score'].isnull()
data[credit_score_nan].head()
```

Out[9]:

	Loan ID	Customer ID	Loan Status	Current Loan Amount	Term	Credit Score	Years in current job	Home Ownership	Age
7	0000afa6-8902-4f8f-b870-25a8fdad0aeb	e49c1a82-a0f7-45e8-9f46-2f75c43f9fbc	Charged Off	24613	Long Term	NaN	6 years	Rent	NaN
8	00011dfc-31c1-4178-932a-fbeb3f341efb	ef6e098c-6c83-4752-8d00-ff793e476b8c	Fully Paid	10036	Short Term	NaN	5 years	Rent	NaN
12	00029f9f-0cc5-4d4e-aabc-ea4a7fe74e12	afbc2fa3-3bad-4d48-b691-829aed78bad5	Charged Off	17980	Short Term	NaN	< 1 year	Own Home	NaN
20	00038a08-f058-4add-a8ed-497b91672a9e	727bb429-dfa1-41c4-a347-23230e23949f	Charged Off	16929	Long Term	NaN	3 years	Home Mortgage	NaN
22	0003b749-307f-4830-9fb4-9db7ed1b1c48	998dc43c-f9ce-466e-bdaa-7057b0bbb9cd	Fully Paid	7228	Short Term	NaN	3 years	Home Mortgage	NaN



*Creating a new dataframe as df that has "Credit Score" values

In [10]:

```
df = data[data['Credit Score'].notnull()]
```

In [11]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 160098 entries, 0 to 256983
Data columns (total 19 columns):
Loan ID                160098 non-null object
Customer ID            160098 non-null object
Loan Status            160098 non-null object
Current Loan Amount    160098 non-null int64
Term                   160098 non-null object
Credit Score           160098 non-null float64
Years in current job    160098 non-null object
Home Ownership          160098 non-null object
Annual Income           160098 non-null float64
Purpose                 160098 non-null object
Monthly Debt           160098 non-null object
Years of Credit History 160098 non-null float64
Months since last delinquent 72761 non-null float64
Number of Open Accounts 160098 non-null int64
Number of Credit Problems 160098 non-null int64
Current Credit Balance  160098 non-null int64
Maximum Open Credit     160098 non-null object
Bankruptcies            159756 non-null float64
Tax Liens               160085 non-null float64
dtypes: float64(6), int64(4), object(9)
memory usage: 24.4+ MB
```

*convert Credit Score column from float to int type

In [12]:

```
df['Credit Score'] = df['Credit Score'].astype(np.int64)
```

```
/anaconda/envs/python2/lib/python2.7/site-packages/ipykernel/__main__.
```

```
py:1: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

(<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

```
if __name__ == '__main__':
```

In [13]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 160098 entries, 0 to 256983
Data columns (total 19 columns):
Loan ID                160098 non-null object
Customer ID           160098 non-null object
Loan Status           160098 non-null object
Current Loan Amount    160098 non-null int64
Term                  160098 non-null object
Credit Score          160098 non-null int64
Years in current job   160098 non-null object
Home Ownership         160098 non-null object
Annual Income         160098 non-null float64
Purpose               160098 non-null object
Monthly Debt          160098 non-null object
Years of Credit History 160098 non-null float64
Months since last delinquent 72761 non-null float64
Number of Open Accounts 160098 non-null int64
Number of Credit Problems 160098 non-null int64
Current Credit Balance  160098 non-null int64
Maximum Open Credit    160098 non-null object
Bankruptcies          159756 non-null float64
Tax Liens              160085 non-null float64
dtypes: float64(5), int64(5), object(9)
memory usage: 24.4+ MB
```

*Dividing "Credit Score" by 10, for the values more than 800

In [14]:

```
df['Credit Score'] = df['Credit Score'].map(lambda x: x/10 if x > 800 else x)
```

```
/anaconda/envs/python2/lib/python2.7/site-packages/ipykernel/__main__.
py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>
(<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

```
if __name__ == '__main__':
```


In [15]:

```
df['Credit Score'][(df['Credit Score'] > 800)].count()
```

Out[15]:

0

*Now, we don't have credit score more than 800

*Convert Monthly Debt from object to float64, removing the \$

In [16]:

```
df['Monthly Debt'] = df['Monthly Debt'].replace( '[\\$, ]', '', regex=True).astype(float)
df['Monthly Debt'].head()
```

/anaconda/envs/python2/lib/python2.7/site-packages/ipykernel/__main__.

py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

(<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

```
if __name__ == '__main__':
```

Out[16]:

0 584.03

1 1106.04

2 1321.85

3 751.92

4 355.18

Name: Monthly Debt, dtype: float64

*Converting "Maximum Open Credit" from object to float64

In [17]:

```
df['Maximum Open Credit'] = df['Maximum Open Credit'].convert_objects(convert_numeric=True)
df['Maximum Open Credit'].head()
```

```
/anaconda/envs/python2/lib/python2.7/site-packages/ipykernel/__main__.py:1: FutureWarning: convert_objects is deprecated. Use the data-type specific converters pd.to_datetime, pd.to_timedelta and pd.to_numeric.
  if __name__ == '__main__':
/anaconda/envs/python2/lib/python2.7/site-packages/ipykernel/__main__.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>
(<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

```
if __name__ == '__main__':
```

Out[17]:

```
0    16056.0
1    19149.0
2    28335.0
3    43915.0
4    37081.0
Name: Maximum Open Credit, dtype: float64
```

*Home Ownership: converting "HaveMortgage" to "Home Mortgage"

In [18]:

```
df['Home Ownership'] = df['Home Ownership'].map(lambda x: 'Home Mortgage' if x == 'HaveMortgage' else x)
```

```
/anaconda/envs/python2/lib/python2.7/site-packages/ipykernel/__main__.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>
(<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

```
if __name__ == '__main__':
```

*Purpose: converting "other" to "Other"

In [19]:

```
#Purpose: other == Other

df['Purpose'] = df['Purpose'].map(lambda x: 'Other' if x == 'other' else x)
```

/anaconda/envs/python2/lib/python2.7/site-packages/ipykernel/__main__.

py:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

(<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

app.launch_new_instance()

*Checking the unique values of "Home Ownership" and "Purpose"

In [20]:

```
print "Home Ownership: ", df['Home Ownership'].unique()
print "Purpose: ", df['Purpose'].unique()
```

Home Ownership: ['Home Mortgage' 'Own Home' 'Rent']

Purpose: ['Debt Consolidation' 'Other' 'Business Loan' 'Home Improvements'

'Buy House' 'Medical Bills' 'Take a Trip' 'Buy a Car'

'Educational Expenses']

*We have missing values in "Month since last delinquent" column. Let's take a look at the relationship between NA values in "Month since last delinquent" column and other "problem columns" such as "Purpose", "Number of Credit Problems", "Bankruptcies", and "Tax Liens"

In [21]:

```
msld_nan = df['Months since last delinquent'].isnull()

print df[msld_nan]['Purpose'].value_counts()

print df[msld_nan]['Number of Credit Problems'].value_counts()

print df[msld_nan]['Bankruptcies'].value_counts()

print df[msld_nan]['Tax Liens'].value_counts()
```

```
Debt Consolidation      69495
Other                   8295
Home Improvements      4513
Business Loan          1728
Buy a Car              1205
Medical Bills          942
Take a Trip            552
Buy House              515
Educational Expenses    92
```

Name: Purpose, dtype: int64

```
0    76539
1     9711
2      808
3     179
4       63
5       21
6       11
7        3
9        1
8        1
```

Name: Number of Credit Problems, dtype: int64

```
0.0    78011
1.0    9016
2.0     266
3.0      35
4.0       5
5.0        2
```

Name: Bankruptcies, dtype: int64

```
0.0    86179
1.0     837
2.0     220
3.0      48
4.0      28
5.0      16
6.0       5
9.0       1
7.0       1
```

Name: Tax Liens, dtype: int64

*Most of NA values in "Month since last delinquent" column have: purpose of Debt Consolidation; 0 Number of Credit Problems; 0 Bankruptcies; 0 Tax Liens

*Now, taking a look at correlations

In [22]:

```
df.corr()
```

Out[22]:

	Current Loan Amount	Credit Score	Annual Income	Monthly Debt	Years of Credit History	Months since last delinquent	Number of Open Accounts
Current Loan Amount	1.000000	-0.233733	0.335476	0.435983	0.152230	-0.044143	0.201689
Credit Score	-0.233733	1.000000	0.009421	-0.095841	0.087433	0.046790	-0.046681
Annual Income	0.335476	0.009421	1.000000	0.452869	0.146897	-0.063535	0.139885
Monthly Debt	0.435983	-0.095841	0.452869	1.000000	0.188655	-0.056744	0.409085
Years of Credit History	0.152230	0.087433	0.146897	0.188655	1.000000	-0.039912	0.130126
Months since last delinquent	-0.044143	0.046790	-0.063535	-0.056744	-0.039912	1.000000	-0.035585
Number of Open Accounts	0.201689	-0.046681	0.139885	0.409085	0.130126	-0.035585	1.000000
Number of Credit Problems	-0.072039	-0.055539	-0.010737	-0.049905	0.061586	0.088823	-0.012141
Current Credit Balance	0.331326	-0.015774	0.283434	0.476020	0.207288	-0.023838	0.231627
Maximum Open Credit	0.038276	0.017475	0.040757	0.037311	0.027981	0.001044	0.030191
Bankruptcies	-0.095954	-0.039888	-0.043599	-0.076680	0.062593	0.111428	-0.020951
Tax Liens	0.014140	-0.028144	0.041296	0.022576	0.020989	0.007285	0.005294

*Let's take a look at "Months since last delinquent" column

In [23]:

```
df['Months since last delinquent'].max()
```

Out[23]:

176.0

*Let's fill in NA in 'Months since last delinquent' with 200 which is even more than max number of credit problems (176), because we noticed above that most NA in 'Months since last delinquent' has "least" problems

In [24]:

```
df["Months since last delinquent"].fillna(value=200, inplace=True)
```

```
//anaconda/envs/python2/lib/python2.7/site-packages/pandas/core/generic.py:3191: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>
(<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)
self._update_inplace(new_data)

In [25]:

```
df['Months since last delinquent'].describe()
```

Out[25]:

```
count      160098.000000  
mean         124.948338  
std           83.539069  
min           0.000000  
25%           35.000000  
50%           200.000000  
75%           200.000000  
max           200.000000  
Name: Months since last delinquent, dtype: float64
```

In [26]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 160098 entries, 0 to 256983
Data columns (total 19 columns):
Loan ID                160098 non-null object
Customer ID           160098 non-null object
Loan Status           160098 non-null object
Current Loan Amount    160098 non-null int64
Term                  160098 non-null object
Credit Score          160098 non-null int64
Years in current job   160098 non-null object
Home Ownership         160098 non-null object
Annual Income          160098 non-null float64
Purpose               160098 non-null object
Monthly Debt          160098 non-null float64
Years of Credit History 160098 non-null float64
Months since last delinquent 160098 non-null float64
Number of Open Accounts 160098 non-null int64
Number of Credit Problems 160098 non-null int64
Current Credit Balance 160098 non-null int64
Maximum Open Credit    160095 non-null float64
Bankruptcies          159756 non-null float64
Tax Liens              160085 non-null float64
dtypes: float64(7), int64(5), object(7)
memory usage: 24.4+ MB
```

*Tax Liens is highly correlated with Number of credit problems = 0.59, so let's only include none missing values in Tax Liens for our new dataframe df

In [27]:

```
df = df[df['Tax Liens'].notnull()]
```

In [28]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 160085 entries, 0 to 256983
Data columns (total 19 columns):
Loan ID                160085 non-null object
Customer ID           160085 non-null object
Loan Status            160085 non-null object
Current Loan Amount    160085 non-null int64
Term                   160085 non-null object
Credit Score           160085 non-null int64
Years in current job   160085 non-null object
Home Ownership         160085 non-null object
Annual Income          160085 non-null float64
Purpose                160085 non-null object
Monthly Debt           160085 non-null float64
Years of Credit History 160085 non-null float64
Months since last delinquent 160085 non-null float64
Number of Open Accounts 160085 non-null int64
Number of Credit Problems 160085 non-null int64
Current Credit Balance  160085 non-null int64
Maximum Open Credit     160082 non-null float64
Bankruptcies           159756 non-null float64
Tax Liens              160085 non-null float64
dtypes: float64(7), int64(5), object(7)
memory usage: 24.4+ MB
```

*Let's only include none missing values in Bankruptcies for our new dataframe df

In [29]:

```
df = df[df['Bankruptcies'].notnull()]
```


In [30]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 159756 entries, 0 to 256983
Data columns (total 19 columns):
Loan ID                159756 non-null object
Customer ID            159756 non-null object
Loan Status            159756 non-null object
Current Loan Amount    159756 non-null int64
Term                   159756 non-null object
Credit Score           159756 non-null int64
Years in current job   159756 non-null object
Home Ownership          159756 non-null object
Annual Income          159756 non-null float64
Purpose                159756 non-null object
Monthly Debt           159756 non-null float64
Years of Credit History 159756 non-null float64
Months since last delinquent 159756 non-null float64
Number of Open Accounts 159756 non-null int64
Number of Credit Problems 159756 non-null int64
Current Credit Balance  159756 non-null int64
Maximum Open Credit     159753 non-null float64
Bankruptcies            159756 non-null float64
Tax Liens               159756 non-null float64
dtypes: float64(7), int64(5), object(7)
memory usage: 24.4+ MB
```

*Taking a look at "Maximum Open Credit" column

In [31]:

```
moc_nan = df['Maximum Open Credit'].isnull()
moc_nan.value_counts()
```

Out[31]:

```
False    159753
True         3
Name: Maximum Open Credit, dtype: int64
```

*Selecting none missing values in Maximum Open Credit for our new dataframe df

In [32]:

```
df = df[df['Maximum Open Credit'].notnull()]
```

In [33]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 159753 entries, 0 to 256983
Data columns (total 19 columns):
Loan ID                159753 non-null object
Customer ID            159753 non-null object
Loan Status            159753 non-null object
Current Loan Amount    159753 non-null int64
Term                  159753 non-null object
Credit Score          159753 non-null int64
Years in current job   159753 non-null object
Home Ownership         159753 non-null object
Annual Income          159753 non-null float64
Purpose               159753 non-null object
Monthly Debt          159753 non-null float64
Years of Credit History 159753 non-null float64
Months since last delinquent 159753 non-null float64
Number of Open Accounts 159753 non-null int64
Number of Credit Problems 159753 non-null int64
Current Credit Balance  159753 non-null int64
Maximum Open Credit    159753 non-null float64
Bankruptcies           159753 non-null float64
Tax Liens              159753 non-null float64
dtypes: float64(7), int64(5), object(7)
memory usage: 24.4+ MB
```

*Now, we don't have any missing values

*Taking look at "Loan ID" and "Customer ID" columns

In [34]:

```
print "Total unique Loan IDs: ", df['Loan ID'].nunique()
print "Total unique Customer IDs: ", df['Customer ID'].nunique()
print "Total entries: ", len(df)
```

```
Total unique Loan IDs:  135696
Total unique Customer IDs:  135696
Total entries:  159753
```

In [35]:

```
df['Loan ID'].value_counts()
```

Out[35]:

```
9d601e9f-2e11-42f6-868a-82ba3aa9e401    4
bbf7a3d6-c415-417e-b6a1-5095a6a5fcca    4
80bf6f56-b3f3-41d8-bb26-191ead5a8a4c    4
```

255cc82a-2733-48bd-a784-aa61a494833b	4
21d50a6a-0e3d-49d9-843d-e4c9e2959e11	4
2f7500e8-695a-4d4d-b91c-bdfaedb06745	4
536c1d67-0e1a-4060-909b-b23b7d6eff50	4
4e800cb9-955b-45f9-85d5-196c3054f2f1	4
f1dc64e8-b329-480a-b289-e793e954fd3c	4
a4bd5c89-275a-4539-9e27-5c256845bd2c	4
943b437f-b6aa-4581-9ca7-611a5fffffc4a	4
16a5ca88-8f09-4a8e-a40c-9e9bc394abc7	4
6e766ede-75e9-4b26-b781-e268faf258f0	4
e6559d6b-42fa-4285-a565-200b693fb976	4
91037789-b0a2-4d49-89dc-78c27b5e5297	4
9f20c0be-fac2-4400-8c2a-fa6063a123a0	4
b1e4e451-4a27-4a35-9d48-3fc55702992b	4
66bb3e99-6468-4193-a7d3-4d0134581e30	4
b8dbe5c6-6c6d-4b26-850f-8f52c4303383	4
55931e3a-d5d2-4be5-94c6-d4c8bb9d4844	4
c7f0e509-28b8-4349-99db-f863fb067237	4
16be66d9-607a-43a4-8fbf-5327a16708c1	4
889bb6f2-91f1-46c6-ba21-92f37ea2e662	4
93663d95-1686-4fb7-8d4f-b06c237d5689	4
8a036c7c-9836-4763-8f69-980b23e3f381	4
14c5fb07-d3bd-4fc8-aad0-54d21f491daf	4
f57c84b6-cf4d-4023-ab5d-8ae55a708524	4
e94aceda-a038-4444-bb52-a481f2ac5472	4
2e03e705-c9eb-474c-9e45-6e898d75dbd8	4
cadb86e8-b6da-420f-95ce-a477f5cb3040	4
..	
ea9580a3-53e8-4ae8-b986-b2ceb009a613	1
838674cc-5d8a-487b-99db-34ed9a417d87	1
0a9c74d7-35f3-4f5c-ae7f-d4963fe25c74	1
81cf9604-580e-4360-a39e-ad42f3bfe840	1
9aa79939-a953-4472-aa07-2b35f04a7100	1
8fedaaaf8-f844-4d8d-8b35-4aeae3862844	1
3d8ab1c0-42eb-4495-9aa8-1afb4fa72b1b	1
0fd0dd72-a6e7-4971-8f19-d7b820b6c84e	1
14ba0515-b8a8-4bce-873b-d6143df45b13	1
fb4aade-5491-48cd-b87c-fdcbb615a771	1
97a385bc-b849-4fc0-997e-bfa9db1878af	1
a7933509-e00c-402c-8211-819d7bcdac06	1
ee215047-6034-484e-b98b-8659c9d6ab50	1
f6a9f7a0-e6ce-4bba-a489-ce54096fdd9c	1
d4f2074f-e816-4724-a984-732a8844b7e9	1
6fb2fadc-92b3-4704-9aae-78236b924fa4	1
17e3da6f-2161-4864-b7b3-682cca6430d2	1
36b424f8-a6b8-443b-a46e-b08b1818ff60	1
b480f024-aeb7-4a1c-96be-fb58f99b2301	1
a28ca379-87b6-4332-a0f5-acdfc2c75701	1
e8b0b33a-ad7c-4dc5-91e8-8df5deb35735	1
f4cf3374-4203-4cb1-8fe6-45108ccebda1	1
d7360a03-c7a9-4574-8ac7-a7fc0cd636a1	1
b6e48ebc-fb87-4c90-9ae5-12d98262f4fc	1
e1d5d2bd-c029-4b55-a0c0-64c0997a677c	1

```
201c3791-b28e-4605-a130-313fe0a2069e    1
d070b35b-cb35-474c-9e72-f26ef01c3bbd    1
2f282acf-636a-456f-9b13-aec30d35b817    1
efc6e397-4743-4843-9477-2a0269abae8a    1
5b11b29a-700e-42fa-9992-14067457ec62    1
Name: Loan ID, dtype: int64
```

*Let's drop the duplicate Loan ID values & keep first row, and set new dataframe df

In [36]:

```
df = df.drop_duplicates(['Loan ID'], keep = 'first')
```

In [37]:

```
print "Total unique Loan IDs: ", df['Loan ID'].nunique()
print "Total unique Customer IDs: ", df['Customer ID'].nunique()
print "Total entries: ", len(df)
```

```
Total unique Loan IDs:  135696
Total unique Customer IDs:  135696
Total entries:  135696
```

In [38]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 135696 entries, 0 to 256983
Data columns (total 19 columns):
Loan ID                135696 non-null object
Customer ID            135696 non-null object
Loan Status            135696 non-null object
Current Loan Amount    135696 non-null int64
Term                  135696 non-null object
Credit Score           135696 non-null int64
Years in current job    135696 non-null object
Home Ownership          135696 non-null object
Annual Income           135696 non-null float64
Purpose                 135696 non-null object
Monthly Debt           135696 non-null float64
Years of Credit History 135696 non-null float64
Months since last delinquent 135696 non-null float64
Number of Open Accounts 135696 non-null int64
Number of Credit Problems 135696 non-null int64
Current Credit Balance  135696 non-null int64
Maximum Open Credit     135696 non-null float64
Bankruptcies            135696 non-null float64
Tax Liens               135696 non-null float64
dtypes: float64(7), int64(5), object(7)
memory usage: 20.7+ MB
```

*Double check that we don't have duplicated rows

In [39]:

```
print "Total unique Customer IDs: ", df['Customer ID'].nunique()  
print "Total entries: ", len(df)
```

Total unique Customer IDs: 135696

Total entries: 135696

*Let's take a look at "Years in current job" column - removing the strings by deleting "years and +"; converting it to int; converting n/a and <1 to 0

In [40]:

```
df['Years in current job'] = df['Years in current job'].map(lambda x: '0' if x == 'n/a' else x)  
df['Years in current job'] = df['Years in current job'].map(lambda x: '0' if x == '<1' else x)  
df['Years in current job'] = df['Years in current job'].replace( '[\+ years]', '', regex=True)
```

In [41]:

```
df['Years in current job'].unique()
```

Out[41]:

```
array([10,  4,  6,  0,  2,  5,  3,  1,  7,  9,  8])
```

*Now, we we have integers only for "Years in current job" column

*Converting "Loan Status" to binary output for easy modeling : Fully paid = 1

In [42]:

```
df['Loan Status'] = df['Loan Status'].map(lambda x: 1 if x == 'Fully Paid' else 0)
```

In [43]:

```
df['Loan Status'].unique()
```

Out[43]:

```
array([1, 0])
```

*Converting "Term" to binary output for easy modeling: Long Term = 1

In [44]:

```
df['Term'] = df['Term'].map(lambda x: 1 if x == 'Long Term' else 0)
```

In [45]:

```
df['Term'].unique()
```

Out[45]:

```
array([0, 1])
```

In [46]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 135696 entries, 0 to 256983
Data columns (total 19 columns):
Loan ID                135696 non-null object
Customer ID            135696 non-null object
Loan Status            135696 non-null int64
Current Loan Amount    135696 non-null int64
Term                  135696 non-null int64
Credit Score          135696 non-null int64
Years in current job   135696 non-null int64
Home Ownership         135696 non-null object
Annual Income          135696 non-null float64
Purpose               135696 non-null object
Monthly Debt          135696 non-null float64
Years of Credit History 135696 non-null float64
Months since last delinquent 135696 non-null float64
Number of Open Accounts 135696 non-null int64
Number of Credit Problems 135696 non-null int64
Current Credit Balance  135696 non-null int64
Maximum Open Credit    135696 non-null float64
Bankruptcies          135696 non-null float64
Tax Liens              135696 non-null float64
dtypes: float64(7), int64(8), object(4)
memory usage: 20.7+ MB
```

Feature engineering 1: Credit Utilization Rate = sum of outstanding balance / credit card's limit; In our case,
Credit Utilization Rate = Current Credit Balance / Maximum Open Credit

In [47]:

```
df['Credit Utilization'] = df['Current Credit Balance'].div(df['Maximum Open Credit'])
```



In [48]:

```
df['Credit Utilization'].describe()
```

Out[48]:

```
count      135696.000000
mean         29.092050
std         1536.194716
min           0.000000
25%          0.369991
50%          0.562006
75%          0.741875
max        236458.000000
Name: Credit Utilization, dtype: float64
```

*Feature engineering 2: Payment Rate = (Monthly Debt * 12) / Annual Income = Annual Payment / Annual Income ; Add 1 to this value*

In [49]:

```
df['Annual Payment'] = df['Monthly Debt']*12
```

In [50]:

```
df['Payment Rate'] = df['Annual Payment'].div(df['Annual Income'] + 1, axis='index')
```

In [51]:

```
df['Payment Rate'].describe()
```

Out[51]:

```
count      135696.000000
mean         0.170673
std         0.079576
min           0.000000
25%          0.110999
50%          0.166998
75%          0.225999
max          0.591982
Name: Payment Rate, dtype: float64
```

*Renaming all the columns to make it simple

In [52]:

```
df = df.rename(columns={'Loan Status': 'y',
                        'Current Loan Amount': 'a',
                        'Term': 'b', 'Credit Score': 'c',
                        'Years in current job': 'd',
                        'Home Ownership': 'e',
                        'Annual Income': 'f',
                        'Purpose': 'g', 'Monthly Debt': 'h',
                        'Years of Credit History': 'i',
                        'Number of Open Accounts': 'j',
                        'Number of Credit Problems': 'k',
                        'Current Credit Balance': 'l',
                        'Maximum Open Credit': 'm',
                        'Bankruptcies': 'n', 'Tax Liens': 'o',
                        'Credit Utilization': 'p', 'Payment Rate': 'q'})

df
```

Out[52]:

	Loan ID	Customer ID	y	a	b	c	d	e	f	g	...	Mo sir de
0	000025bb-5694-4cff-b17d-192b1a98ba44	5ebc8bb1-5eb9-4404-b11b-a6eebc401a19	1	11520	0	741	10	Home Mortgage	33694.0	Debt Consolidation	...	41
1	00002c49-3a29-4bd4-8f67-c8f8fbc1048c	927b388d-2e01-423f-a8dc-f7e42d668f46	1	3441	0	734	4	Home Mortgage	42269.0	Other	...	20
	00002d89-	defce609-										

Data Visualization

In [53]:

```
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

*Plotting "Current Loan Amount" for histogram and boxplot


In [54]:

```
fig = plt.figure()
ax1 = fig.add_subplot(121)
ax2 = fig.add_subplot(122)

ax1.hist(df.a)
ax2.boxplot(df.a)
```

Out[54]:

```
{'boxes': [<matplotlib.lines.Line2D at 0x11aa67d90>],
'caps': [<matplotlib.lines.Line2D at 0x10e755090>,
<matplotlib.lines.Line2D at 0x10e7556d0>],
'fliers': [<matplotlib.lines.Line2D at 0x10e760390>],
'means': [],
'medians': [<matplotlib.lines.Line2D at 0x10e755d10>],
'whiskers': [<matplotlib.lines.Line2D at 0x11aa67e90>,
<matplotlib.lines.Line2D at 0x119f18610>]}
```




*Normalizing it by applying log into "Current Loan Amount"

In [55]:

```
plt.hist(np.log(df.a))
```

Out[55]:

```
(array([ 880., 1740., 2525., 6426., 14170., 17367., 31075.,  
        26720., 22926., 11867.]),  
 array([ 6.88243747,  7.25210188,  7.62176629,  7.9914307 ,  
        8.36109511,  8.73075952,  9.10042393,  9.47008834,  
        9.83975275, 10.20941716, 10.57908157]),  
<a list of 10 Patch objects>)
```



*Now we have the histogram that is skewed to the left

*Now we want to plot x as "Credit Utilization" & y as "Loan Status"


In [56]:

```
x = df['p']
y = df['y']

# make the plot
fig = plt.figure()
ax = fig.add_subplot(111)
ax.scatter(x, y)
ax.set_ylabel('Loan Status', fontsize=20)
ax.set_xlabel('Credit Utilization', fontsize=20)
```

Out[56]:

<matplotlib.text.Text at 0x10e7c0290>



*Note that Loan Status = 1 is Fully Paid and it's more dense on the left; We can conclude that customers who have Loan Status of Fully Paid tend to have lower credit utilization rate

*Now we want to create the dummy variables for categorical values using dmatrices

In [57]:

```
from patsy import dmatrices
y, X = dmatrices('y ~ a + b + c + d + C(e) + f + C(g) + h + i + j + k + l + m + n +
```

In [58]:

```
X.head()
```

Out[58]:

	Intercept	C(e) [T.Own Home]	C(e) [T.Rent]	C(g) [T.Buy House]	C(g) [T.Buy a Car]	C(g)[T.Debt Consolidation]	C(g) [T.Educational Expenses]	C(g)[T.Home Improvement]
0	1.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0
1	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2	1.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0
3	1.0	1.0	0.0	0.0	0.0	1.0	0.0	0.0
4	1.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0

5 rows × 26 columns

In [59]:

```
y.head()
```

Out[59]:

	y
0	1.0
1	1.0
2	1.0
3	1.0
4	1.0

In [60]:

```
X.shape, y.shape
```

Out[60]:

((135696, 26), (135696, 1))

Modeling

In [61]:

```
from sklearn.cross_validation import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.7, random_state=0)
```

```
//anaconda/envs/python2/lib/python2.7/site-packages/sklearn/cross_validation.py:44: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model_selection module into which all the refactored classes and functions are moved. Also note that the interface of the new CV iterators are different from that of this module. This module will be removed in 0.20.
```

```
"This module will be removed in 0.20.", DeprecationWarning)
```

1. Logistic Regression

In [62]:

```
from sklearn.linear_model import LogisticRegression
```

```
logreg = LogisticRegression()
logreg.fit(X_train, y_train)
logreg.score(X_test, y_test)
```

```
//anaconda/envs/python2/lib/python2.7/site-packages/sklearn/utils/validation.py:526: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, 1), for example using ravel().
```

```
y = column_or_1d(y, warn=True)
```

Out[62]:

```
0.73177921344174501
```

In [63]:

```
y_pred_loanstatus = logreg.predict(X_test)
```

In [64]:

```
from sklearn import metrics
print metrics.accuracy_score(y_pred_loanstatus, y_test)
```

```
0.731779213442
```

In [65]:

```
report = metrics.classification_report(y_pred_loanstatus, y_test)
print report
```

	precision	recall	f1-score	support
0.0	0.01	0.51	0.01	140
1.0	1.00	0.73	0.84	40569
avg / total	0.99	0.73	0.84	40709

*Logistic Regression: accuracy score = 0.73; f1 score = 0.84

*Drop Intercept column to fit different models

In [66]:

```
X = X.drop(['Intercept'], axis = 1)
```

In [67]:

```
from sklearn.cross_validation import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size = 0.7, random_s
```

2.. Random Forest Classifier

*Let's take a look at shapes of X, y, X_train, y_train, X_test, y_test

In [68]:

```
X.shape, y.shape
```

Out[68]:

```
((135696, 25), (135696, 1))
```

In [69]:

```
X_train.shape, y_train.shape
```

Out[69]:

```
((94987, 25), (94987, 1))
```

In [70]:

```
X_test.shape, y_test.shape
```

Out[70]:

```
((40709, 25), (40709, 1))
```

*1st method: GridSearchCV

*Converting y_train to array

In [71]:

```
y_train = np.ravel(y_train)
```

In [72]:

```
y_train.shape
```

Out[72]:

```
(94987,)
```

In [73]:

```
from sklearn.grid_search import GridSearchCV
from sklearn.ensemble import RandomForestClassifier

rfc = RandomForestClassifier()

param_grid = {"n_estimators": [100, 200, 300],
              "max_features": [3, 5],
              "max_depth": [10, 20],
              "min_samples_split": [2, 4]}

grid_search = GridSearchCV(estimator=rfc, param_grid=param_grid, n_jobs=-1, cv=2)
grid_search.fit(X_train, y_train)
print grid_search.best_params_
print grid_search.best_score_
print grid_search.best_estimator_
```

```
//anaconda/envs/python2/lib/python2.7/site-packages/sklearn/grid_search
h.py:43: DeprecationWarning: This module was deprecated in version 0.1
8 in favor of the model_selection module into which all the refactored
classes and functions are moved. This module will be removed in 0.20.
  DeprecationWarning)
```

```
{'max_features': 5, 'min_samples_split': 2, 'n_estimators': 200, 'max_
depth': 10}
0.735269036816
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='g
ini',
                      max_depth=10, max_features=5, max_leaf_nodes=None,
                      min_impurity_split=1e-07, min_samples_leaf=1,
                      min_samples_split=2, min_weight_fraction_leaf=0.0,
                      n_estimators=200, n_jobs=1, oob_score=False, random_state=
None,
                      verbose=0, warm_start=False)
```

Now, we can put "best" parameters into random forest classifier; Fitting in training, scoring in testing

In [74]:

```
rfc = RandomForestClassifier(max_features= 5, min_samples_split= 2, n_estimators= 100)

rfc.fit(X_train, y_train)
rfc.score(X_test, y_test)
```

Out[74]:

0.73988552899850157

*Predicting in X_test

In [75]:

```
y_pred_loanstatus= rfc.predict(X_test)
```

In [76]:

```
from sklearn import metrics
print metrics.accuracy_score(y_pred_loanstatus, y_test)
```

0.739885528999

In [77]:

```
report = metrics.classification_report(y_pred_loanstatus, y_test)
print report
```

	precision	recall	f1-score	support
0.0	0.13	0.57	0.21	2480
1.0	0.96	0.75	0.84	38229
avg / total	0.91	0.74	0.81	40709

*Logistic Regression (1st method): accuracy score = 0.74; f1 score = 0.81

*2nd method: GridSearchCV using make_classification; use n_samples = total rows from df (135696), #n_features as number of total columns (26)

In [96]:

```
from sklearn.grid_search import GridSearchCV
from sklearn.datasets import make_classification
from sklearn.ensemble import RandomForestClassifier

# Build a classification task using 3 informative features
X, y = make_classification(n_samples=135696,
                           #n_features=26,
                           #n_informative=3,
                           #n_redundant=0,
                           #n_repeated=0,
                           #n_classes=2,
                           random_state=0,
                           shuffle=False)

rfc = RandomForestClassifier(n_jobs=-1, max_features= 'sqrt' , n_estimators=500, oob_score=True)

param_grid = {'n_estimators': [100, 200, 300],
              'max_features': ['auto', 'sqrt', 'log2']}

grid = GridSearchCV(estimator=rfc, param_grid=param_grid, cv= 2)
grid.fit(X_train, y_train)
print grid.best_params_
print grid.best_score_
print grid.best_estimator_
```

```
{'max_features': 'log2', 'n_estimators': 300}
0.733237179825
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                        max_depth=None, max_features='log2', max_leaf_nodes=None,
                        min_impurity_split=1e-07, min_samples_leaf=1,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        n_estimators=300, n_jobs=-1, oob_score=True, random_state=
None,
                        verbose=0, warm_start=False)
```

*Now, we can put "best" parameters into random forest classifier; Fitting in training, scoring in testing

In [98]:

```
rfc = RandomForestClassifier(max_features= 'log2', n_estimators= 300)
```

In [100]:

```
rfc.fit(X_train, y_train)
rfc.score(X_test, y_test)
```

Out[100]:

0.73887838070205603

*Predicting in X_test

In [101]:

```
y_pred_loanstatus = rfc.predict(X_test)
```

In [102]:

```
from sklearn import metrics
print metrics.accuracy_score(y_pred_loanstatus, y_test)
```

0.738878380702

In [103]:

```
report = metrics.classification_report(y_pred_loanstatus, y_test)
print report
```

	precision	recall	f1-score	support
0.0	0.18	0.54	0.27	3699
1.0	0.94	0.76	0.84	37010
avg / total	0.87	0.74	0.79	40709

*Logistic Regression (2nd method): accuracy score = 0.74; f1 score = 0.79

3.. Gradient boosting - 1st method

In [90]:

```
y_train.shape
```

Out[90]:

(94987,)

*We only need to run GridSearchCV once, so we don't have to re-run here, but we can add "learning rate" parameter for Gradient boosting

In [91]:

```
from sklearn.ensemble import GradientBoostingClassifier

gbc = GradientBoostingClassifier()

param_grid = {"learning_rate": [0.1, 0.5]}

grid_search = GridSearchCV(estimator=gbc, param_grid=param_grid, n_jobs=-1, cv=2)
grid_search.fit(X_train, y_train)
print grid_search.best_params_
print grid_search.best_score_
print grid_search.best_estimator_
```

```
{'learning_rate': 0.1}
0.736279701433
GradientBoostingClassifier(criterion='friedman_mse', init=None,
                           learning_rate=0.1, loss='deviance', max_depth=3,
                           max_features=None, max_leaf_nodes=None,
                           min_impurity_split=1e-07, min_samples_leaf=1,
                           min_samples_split=2, min_weight_fraction_leaf=0.0,
                           n_estimators=100, presort='auto', random_state=None,
                           subsample=1.0, verbose=0, warm_start=False)
```

*Now, we can put "best" parameters(from previous modeling part) & adding learning rate parameter into gradient boosting classifier; Fitting in training, scoring in testing

In [92]:

```
gbc = GradientBoostingClassifier(max_features= 5, min_samples_split= 2, n_estimators= 100)

gbc.fit(X_train, y_train)
gbc.score(X_test, y_test)
```

Out[92]:

```
0.73826426588715022
```

*Predicting in X_test

In [93]:

```
y_pred_loanstatus= gbc.predict(X_test)
```

In [94]:

```
from sklearn import metrics
print metrics.accuracy_score(y_pred_loanstatus, y_test)
```

0.738264265887

In [95]:

```
report = metrics.classification_report(y_pred_loanstatus, y_test)
print report
```


	precision	recall	f1-score	support
0.0	0.21	0.53	0.30	4374
1.0	0.93	0.76	0.84	36335
avg / total	0.85	0.74	0.78	40709

*Gradient Boosting(1st method): accuracy score = 0.74; f1 score = 0.78

Gradient boosting - 2nd method:

In [104]:

```
gbc = GradientBoostingClassifier(max_features= 'log2', n_estimators= 300, learning_
```



In [105]:

```
gbc.fit(X_train, y_train)
gbc.score(X_test, y_test)
```

Out[105]:

0.74212090692475863

In [106]:

```
y_pred_loanstatus= gbc.predict(X_test)
```

In [107]:

```
from sklearn import metrics
print metrics.accuracy_score(y_pred_loanstatus, y_test)
```

0.742120906925

In [108]:

```
report = metrics.classification_report(y_pred_loanstatus, y_test)
print report
```

	precision	recall	f1-score	support
0.0	0.19	0.56	0.28	3727
1.0	0.94	0.76	0.84	36982
avg / total	0.88	0.74	0.79	40709

*Gradient Boosting(2nd method): accuracy score = 0.74; f1 score = 0.79